

Event Queue Dialect:

Bridging Structure and Control

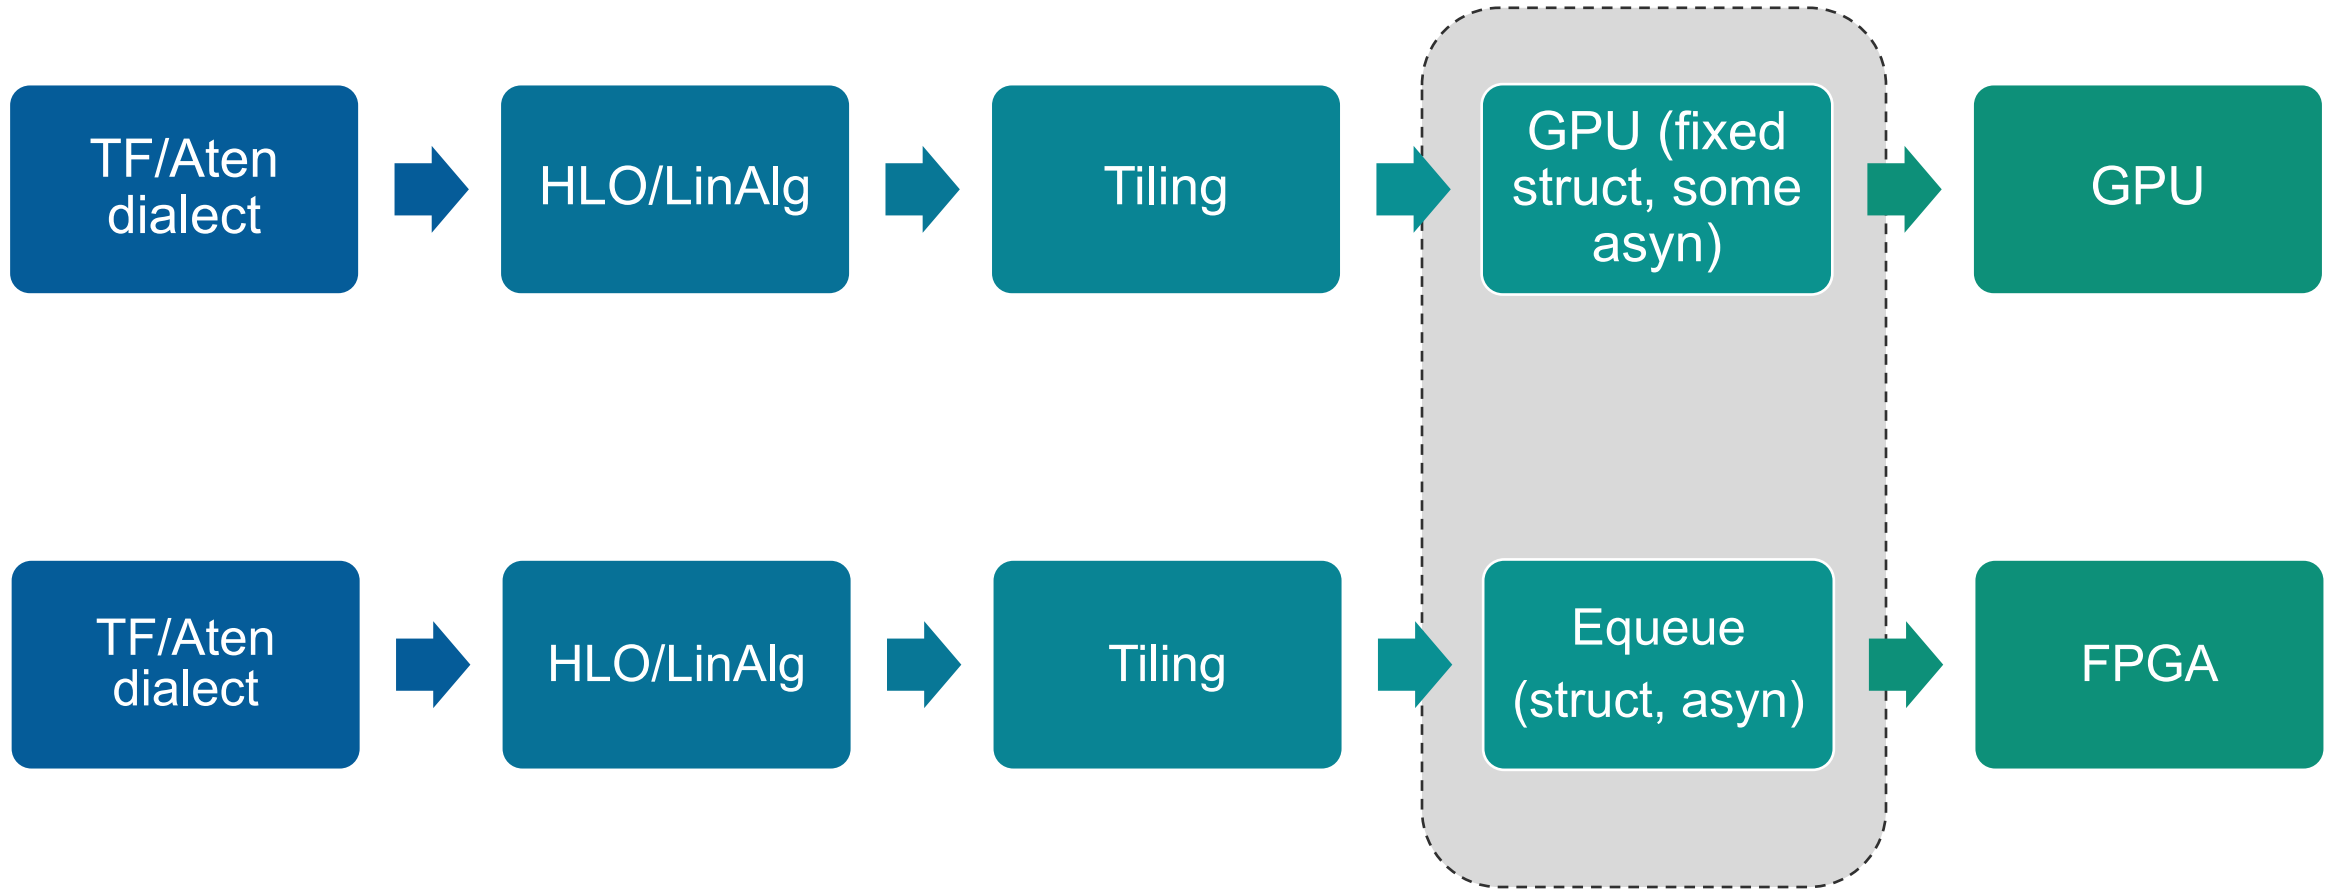
Lowering Pipeline



Lowering Pipeline




Lowering Pipeline



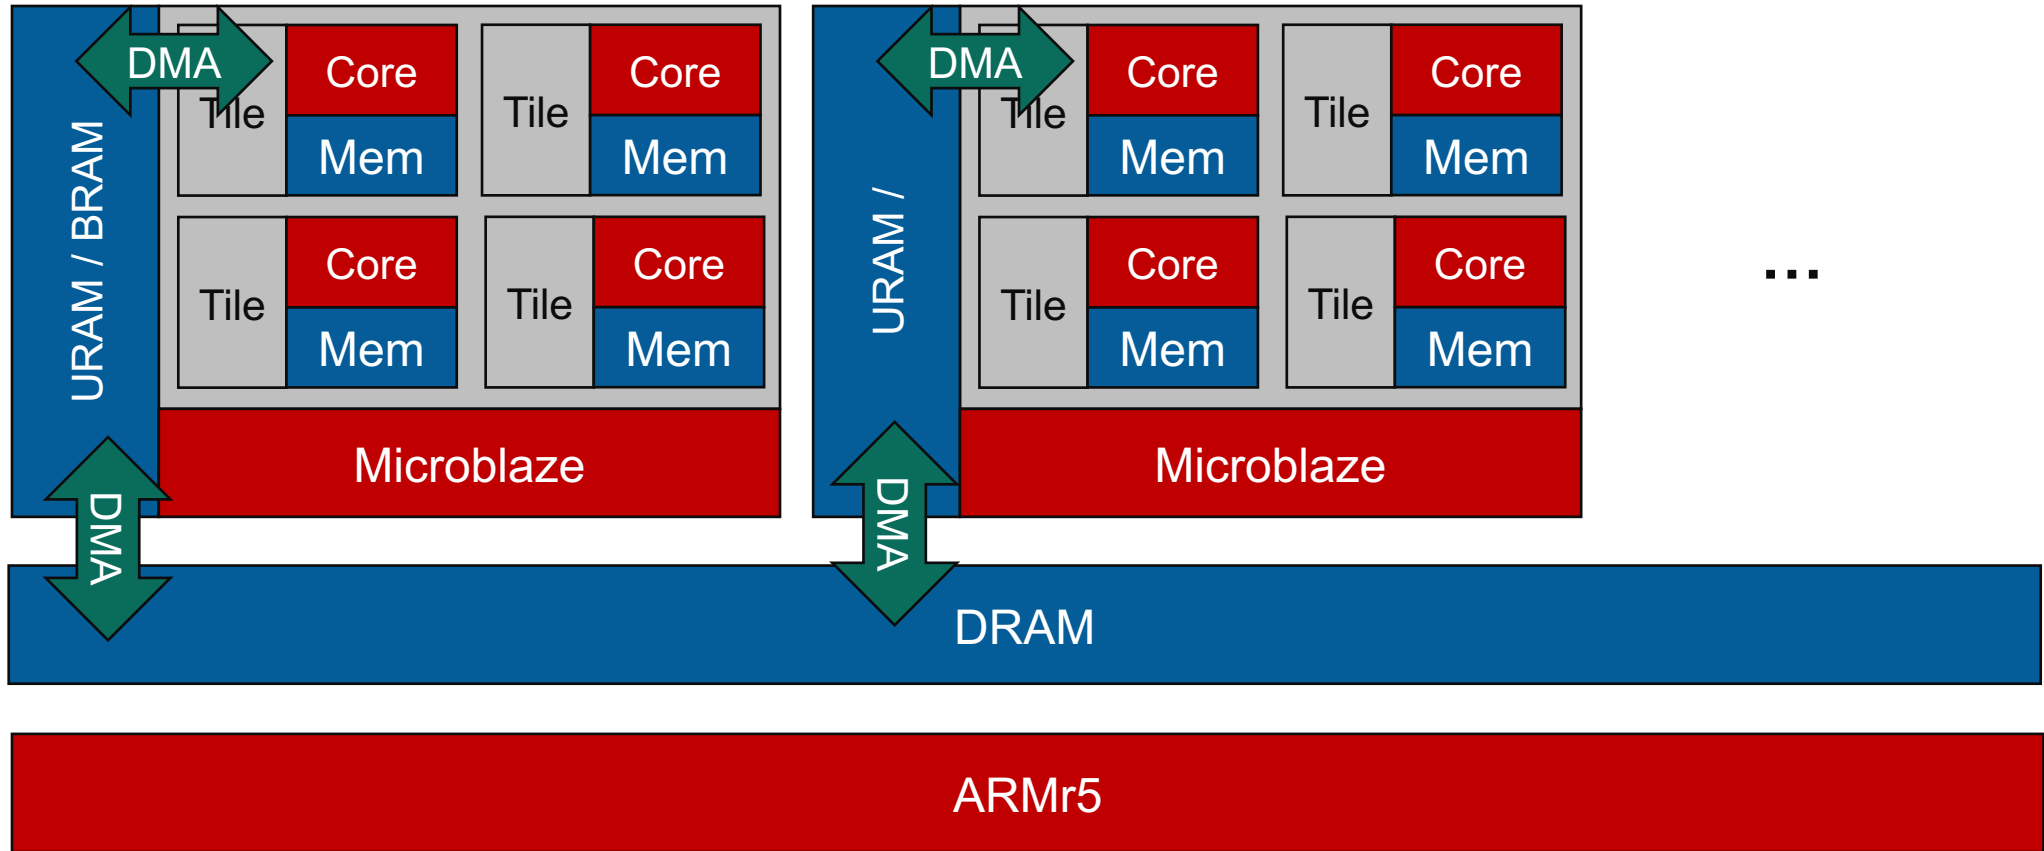
Async Dialect

```
async.func @f(%arg0: tensor<32x8xf32>, %arg1: tensor<8x32xf32>, %arg2: tensor<8x32xf32>, %arg3: tensor<32xf32>) ->
(tensor<32x32xf32>, !async.token) {
  %0:2 = "async.execute"() ( {
    %3 = "some_computation" (%arg0, %arg1) : (tensor<32x8xf32>, tensor<8x32xf32>) -> tensor<32x32xf32>
    "async.yield"(%3) : (tensor<32x32xf32>) -> ()
  }) : () -> (!async.token, tensor<32x32xf32>)
  %1:2 = "async.execute"() ( {
    %3 = "some_computation" (%arg0, %arg1) : (tensor<32x8xf32>, tensor<8x32xf32>) -> tensor<32x32xf32>
    "async.yield"(%3) : (tensor<32x32xf32>) -> ()
  }) : () -> (!async.token, tensor<32x32xf32>)
  %2:2 = "async.execute"(%0#0, %1#0) ( {
    %3 = "some_computation"(%0#1, %1#1, %arg3) : tensor<32x32xf32>, tensor<32x32xf32>, tensor<32xf32> -> tensor<32x32xf32>
    "async.yield"(%3) : (tensor<32x32xf32>) -> ()
  }) : (!async.token, !async.token) -> (!async.token, tensor<32x32xf32>)
  "async.yield"(%2#1, %2#2) : (tensor<32x32xf32>, !async.token) -> ()
}
```



<https://llvm.discourse.group/t/rfc-new-dialect-for-modelling-asynchronous-execution-at-a-higher-level/1345>

Inside Accelerator



Motivation

- ▶ Existing dialects cannot describe programs at a **detailed** enough level
 - **Hardware hierarchy** and **properties**
 - Model **data movement** and **buffer allocation**
 - **Heterogenous processors** with distributed event-based control

Motivation


- ▶ Existing dialects cannot describe programs at a **detailed** enough level
 - **Hardware hierarchy** and **properties**
 - Model **data movement** and **buffer allocation**
 - **Heterogenous processors** with distributed event-based control
- ▶ **Accurate** performance estimation of heterogenous system
 - **High** abstraction level
 - **Fast** estimation
 - **Custom** lowering

High Level Dialect – PyTorch to ATen


```
func @graph( %arg0: tensor<16xf32>, %arg1: tensor<16xf32>) -> tensor<16xf32> {  
    %0 = aten.constant()  
    %arg2 = aten.add(%arg0, %arg1, %0) {acdc_layer_name = "L0-add-0"}  
    return %arg2  
}
```

High Level Dialect – Extending ATen

```
func @graph( %arg0: memref<16xf32>, %arg1:  
    memref<16xf32>) -> memref<16xf32> {  
  
    %m0 = aten.acap_alloc()  
    %s0 = aten.acap_copy(%m0, %arg0)  
    %v0 = aten.acap_tensor_load(%m0, %s0)  
  
    %m1 = aten.acap_alloc()  
    %s1 = aten.acap_copy(%m1, %arg0)  
    %v1 = aten.acap_tensor_load(%m1, %s1)  
  
    %s2 = aten.constant()  
    %m2 = aten.acap_alloc()  
    %v2 = aten.add(%v0, %v1, %s2)  
    %s3 = aten.acap_tensor_store(%v2, %m2)  
    %m3 = aten.acap_alloc()  
    %s4 = aten.acap_copy(%m2, %m3, %s3)  
    %s5 = aten.acap_wait_all(%s4)  
    return %m3  
}
```



Inside Accelerator



EQueue Dialect

- ▶ Model **structure** and capture **hardware properties**
 - ▶ **Explicit** data movements
 - ▶ **Concurrency** between controllers
 - ▶ **Simulator**
-
- ▶ Case study on **input stationary** dataflow on **systolic array**
 - ▶ **Different levels of abstractions** from linalg to equeue dialect

Modeling Hardware

Modeling Hardware

```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```




Memory
(SRAM)

Modeling Hardware

```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```

```
%aie_core = equeue.create_proc "core", AIEngine
```




Modeling Hardware

```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```

```
%aie_core = equeue.create_proc "core", AIEngine
```

```
%dma = equeue.create_dma "dma":()->i32
```




Arbitrary Hardware Hierarchy

```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```

```
%aie_core = equeue.create_proc "core", AIEngine
```

```
%dma = equeue.create_dma "dma":()->i32
```

```
%aie = equeue.create_comp ("aie", %aie_mem, %aie_core, %dma) :  
                                (i32, i32, i32)->i32
```



Arbitrary Hardware Hierarchy


```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```

```
%aie_core = equeue.create_proc "core", AIEngine
```

```
%dma = equeue.create_dma "dma":()->i32
```

```
%aie = equeue.create_comp ("aie", %aie_mem, %aie_core, %dma) :  
                                (i32, i32, i32)->i32
```

```
%dma = equeue.get_comp %aie, "dma"
```




Arbitrary Hardware Hierarchy with FPGA Property

```
%aie_mem = equeue.create_mem "mem", [11], f32, SRAM
```


```
%aie_core = equeue.create_proc "core", AIEngine
```

```
%dma = equeue.create_dma "dma":()->i32
```

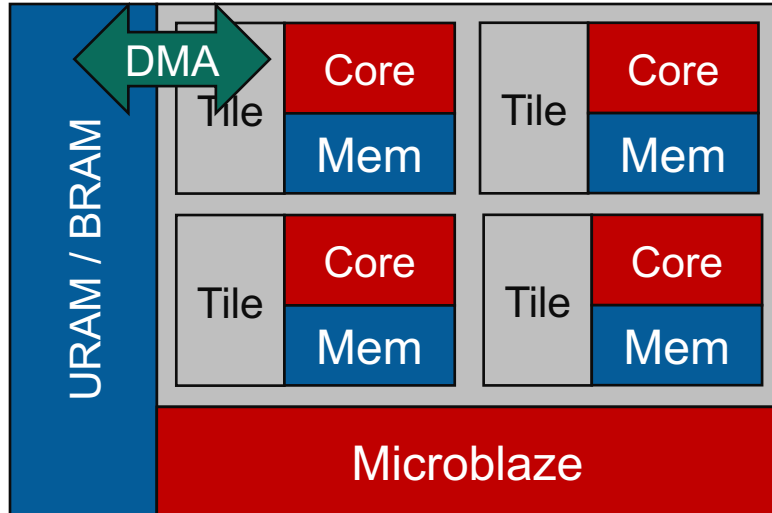
```
%aie = equeue.create_comp ("aie", %aie_mem, %aie_core, %dma) :  
                                (i32, i32, i32)->i32
```



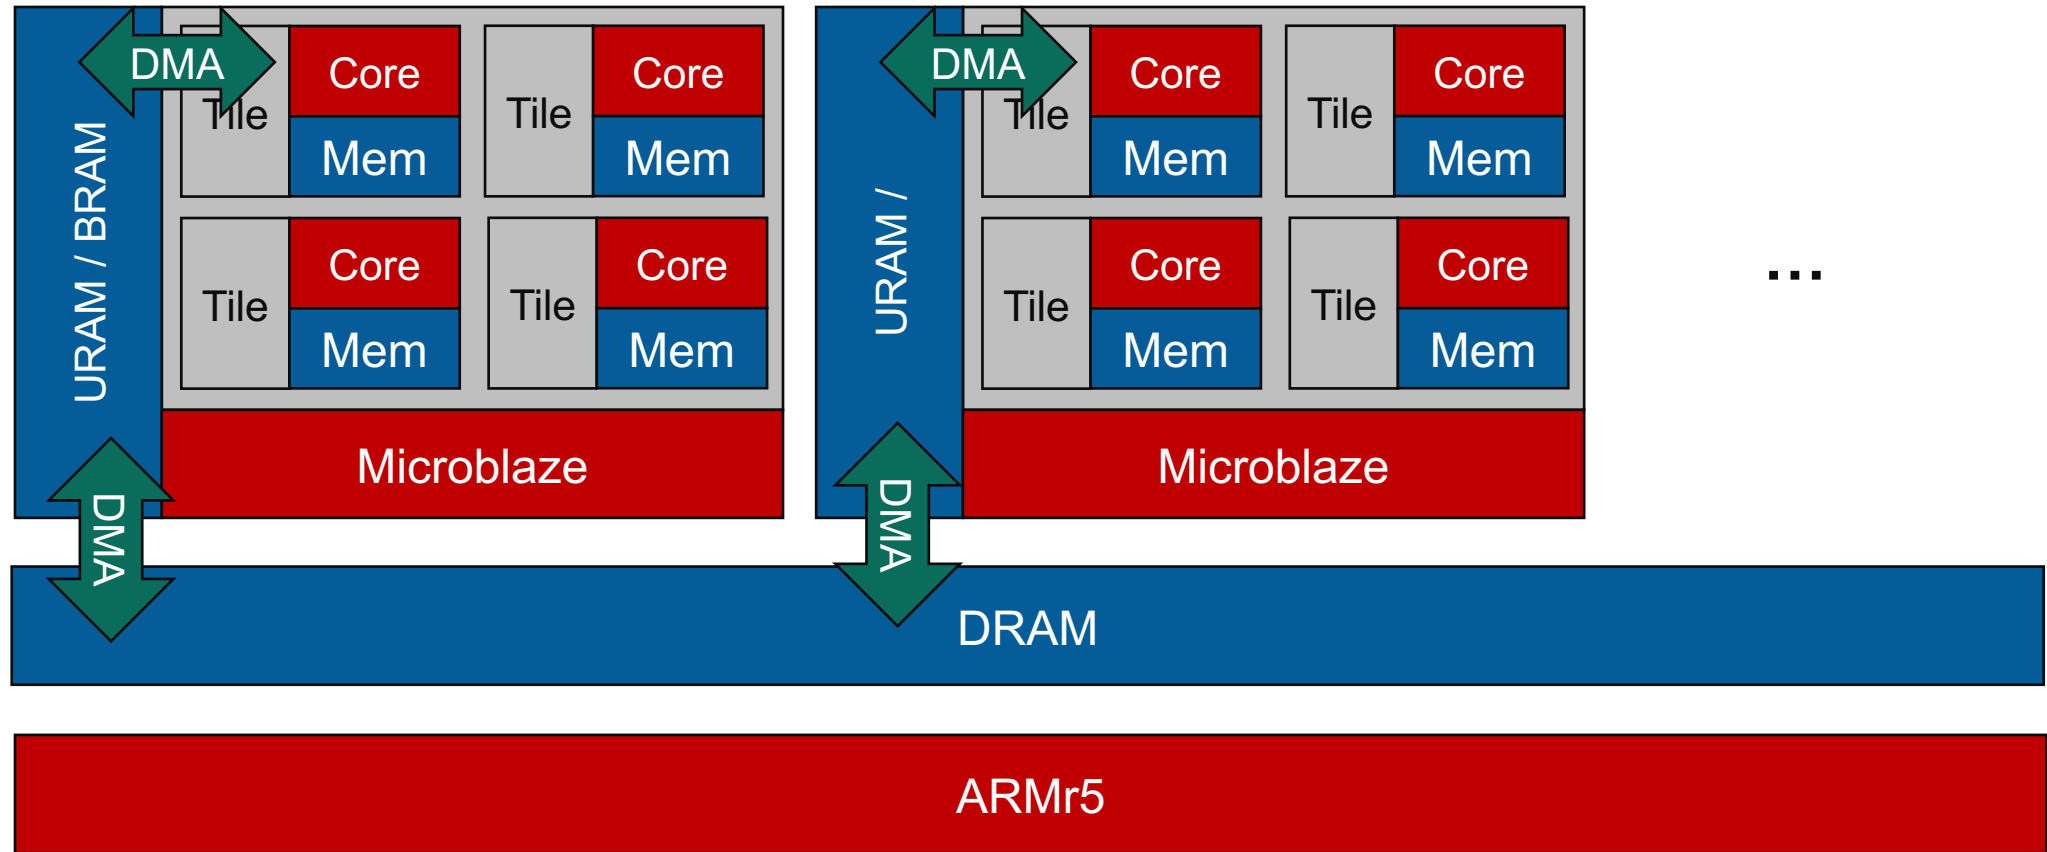
Modeling Large Scale Hardware



Modeling Large Scale Hardware



Modeling Large Scale Hardware



Explicit Data Movements

Explicit Data Movements among Memory

```
%sram = equeue.create_mem "sram", [64], f32, SRAM
```

```
%mem = equeue.create_mem "mem", [4], f32, SRAM
```

```
%sram_buffer = equeue.alloc %sram, [4], f32
```


```
%pe_buffer = equeue.alloc %mem, [4], f32
```

```
%data = equeue.read( %sram_buffer )
```

```
equeue.write( %data, %pe_buffer )
```

```
equeue.dealloc(%sram_buffer)
```

```
equeue.dealloc(%pe_buffer)
```




Concurrency between Controllers

Concurrency

- ▶ A **processor** can issue and receive **launch** operation

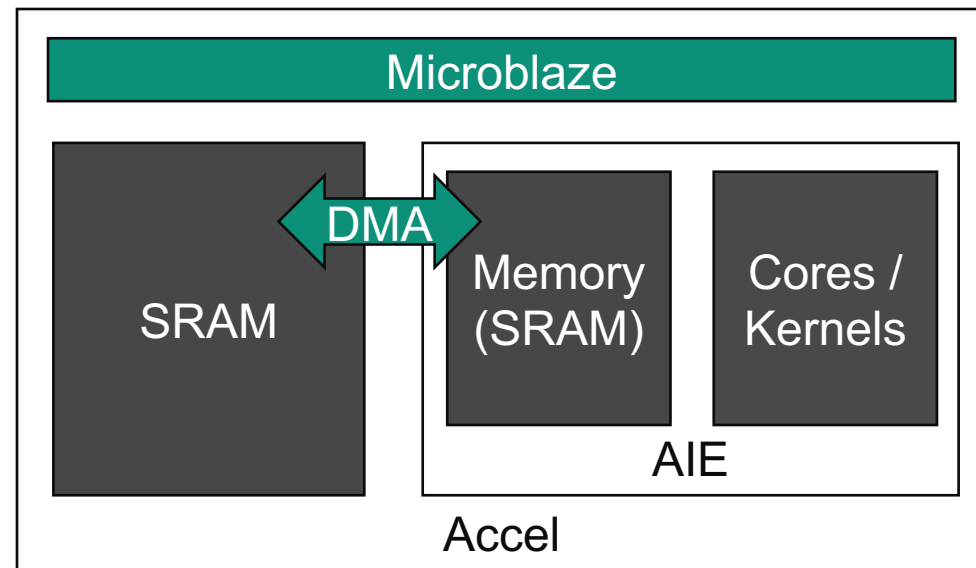
```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    %data = equeue.read(%sram_buffer)
    equeue.write(%data, %aie_buffer)
    ...
    equeue.return
}
```



Concurrency

- ▶ A **processor** can issue and receive **launch** operation

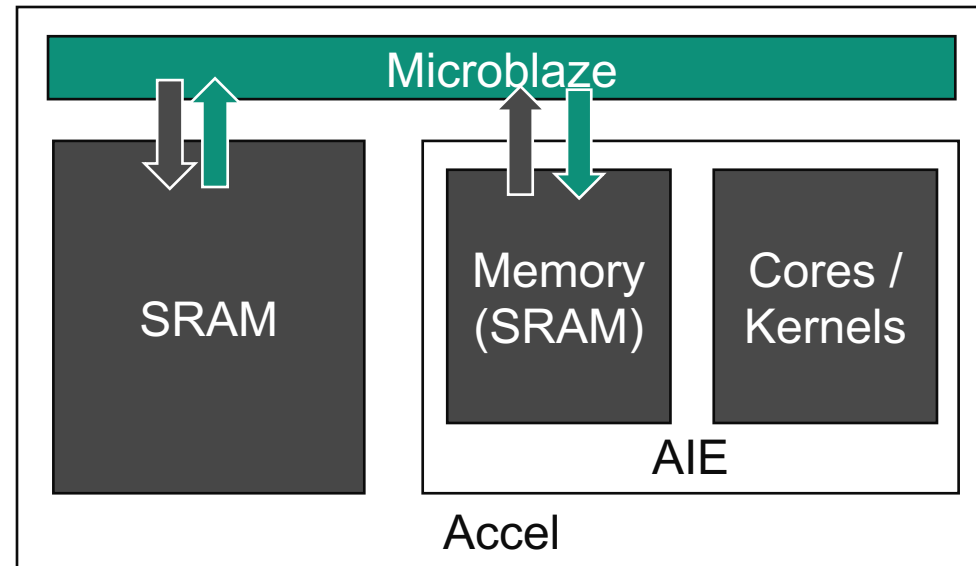
```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```



Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.

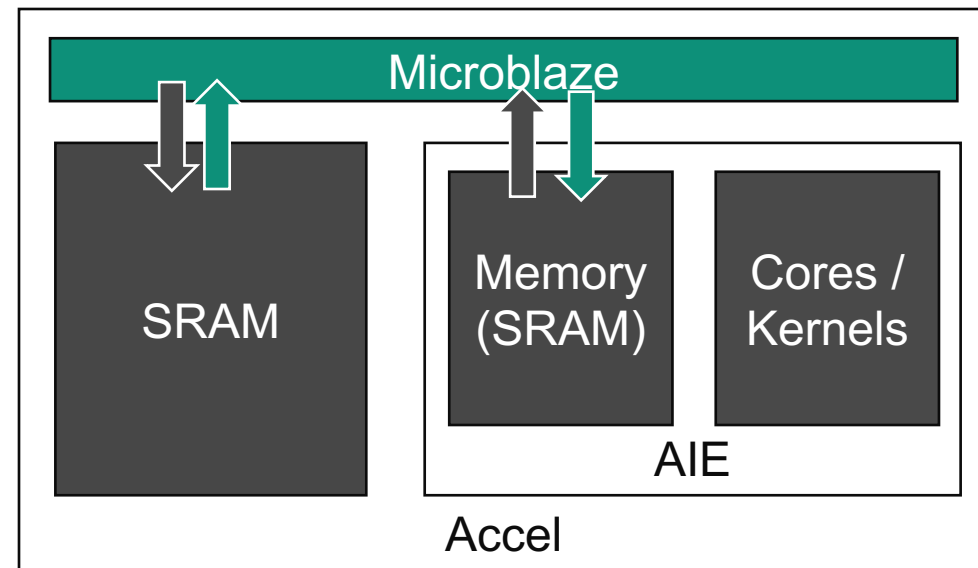
```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```



Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.

```
%microblaze = equeue.create_proc Microblaze
%start = memcpy(%start_memcpy, ...)
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```



Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events


```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```

DMA

memcpy

Microblaze

launch




Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **Sequential Execution**

DMA

```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
➔ %sram_buffer = equeue.alloc %sram, [4], f32
  %aie_buffer = equeue.alloc %mem, [4], f32
  ...
  %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
  %dma)
  ...
  equeue.return
```

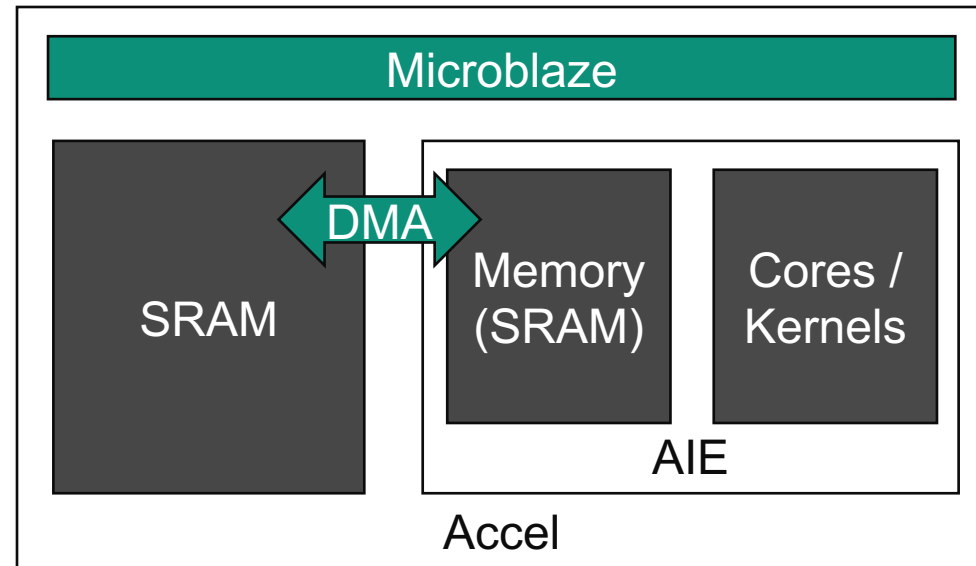


Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **Sequential Execution**

DMA

```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    ➔ %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```

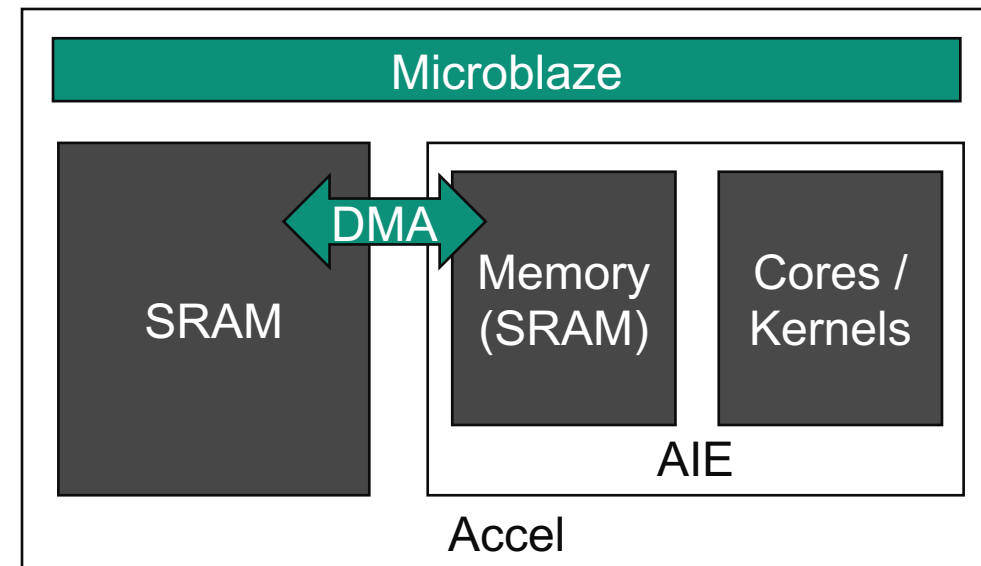



Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **Sequential Execution**

```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    → %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
```

```
}
```



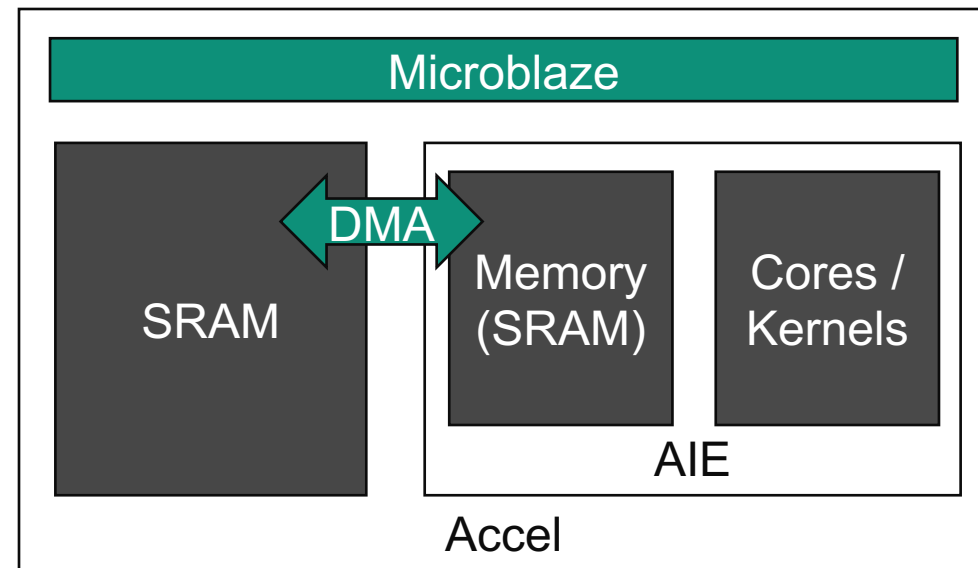
Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **Sequential Execution**

```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
}
```

DMA

memcpy



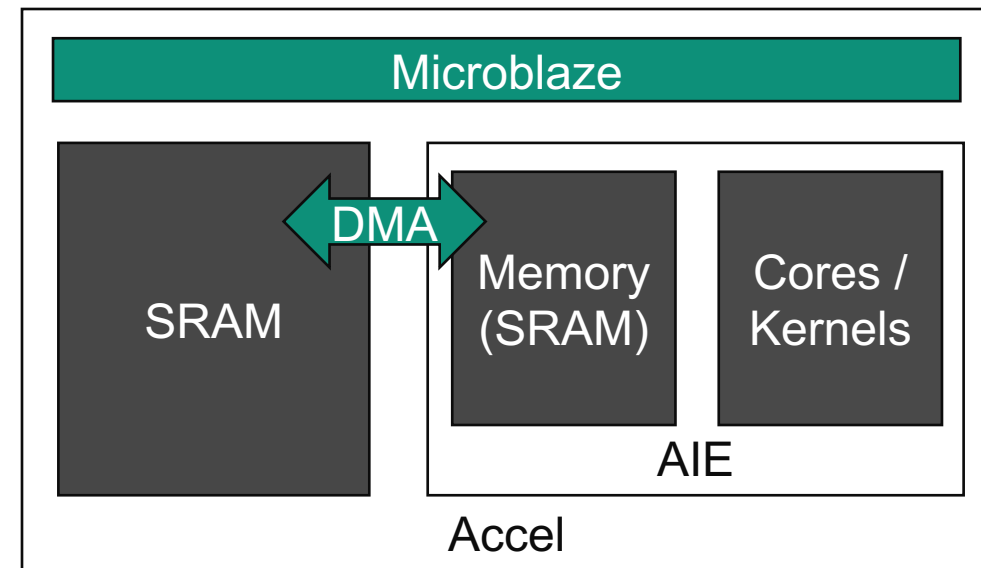
Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **Sequential Execution**

```
%microblaze = equeue.create_proc Microblaze
%done = equeue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = equeue.alloc %sram, [4], f32
    %aie_buffer = equeue.alloc %mem, [4], f32
    ...
    %done_dma = equeue.memcpy(%start_dma, %sram_buffer, %aie_buffer,
    %dma)
    ...
    equeue.return
}
```

DMA

memcpy

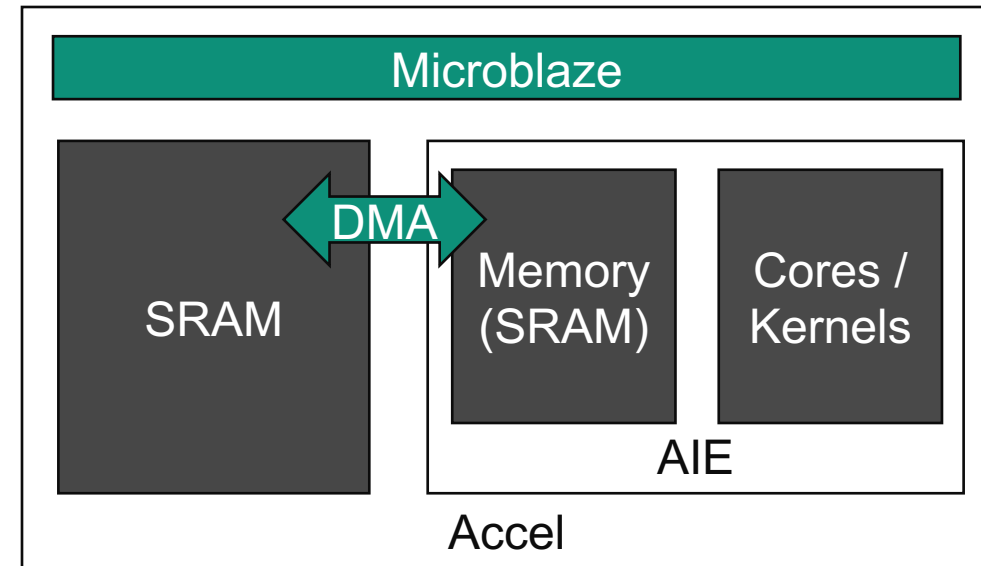


Concurrency

- ▶ A **processor** can issue and receive **launch** operation
 - For launching operation, all **resources** requires explicit handler to **pass control**.
 - Processors are asynchronous: starts on receiving the **start event** ends with a **done event**.
 - **Event queue**: stores operation waiting for events
 - **enqueue.await**: block till events happen

DMA

```
%microblaze = enqueue.create_proc Microblaze
%done = enqueue.launch (%sram, %mem, %dma
= %0, %1, %2) in ( %start, %microblaze){
    %sram_buffer = enqueue.alloc %sram, [4], f32
    %aie_buffer = enqueue.alloc %mem, [4], f32
    ...
    %done_dma = enqueue.memcpy(%sram_buffer, %aie_buffer, %start_dma)
    ...
    enqueue.await(%done_dma)
    enqueue.return
```



A Complete Flow

Rewriting Aten Dialect

```
func @graph( %arg0: memref<16xf32>, %arg1:
    memref<16xf32>) -> memref<16xf32> {

    %m0 = aten.acap_alloc()
    %s0 = aten.acap_copy(%m0, %arg0)
    %v0 = aten.acap_tensor_load(%m0, %s0)

    %m1 = aten.acap_alloc()
    %s1 = aten.acap_copy(%m1, %arg0)
    %v1 = aten.acap_tensor_load(%m1, %s1)

    %s2 = aten.constant()
    %m2 = aten.acap_alloc()
    %v2 = aten.add(%v0, %v1, %s2)
    %s3 = aten.acap_tensor_store(%v2, %m2)
    %m3 = aten.acap_alloc()
    %s4 = aten.acap_copy(%m2, %m3, %s3)
    %s5 = aten.acap_wait_all(%s4)
    return %m3

}
```

```
queue.launch(%kernel, %mem, %dma, %arg0, %arg1, %arg2 =
    %aie, %SRAM, %DMA, %0, %1, %2) in (%s2, %microblaze){
    %start = queue.control_start ()
    %m0 = queue.alloc %mem, [16], f32
    %m1 = queue.alloc %mem, [16], f32
    %m2 = queue.alloc %mem, [16], f32
    %s0 = queue.memcpy(%start, %arg0, %m0, %dma)
    %s1 = queue.memcpy(%start, %arg1, %m1, %dma)
    %s2 = queue.control_and(%s0, %s1)
    %s3= queue.launch(%k0, %k1, %k2 = %m0, %m1, %m2) in
    (%s2, %kernel) {
        %v0 = queue.read(%k0)
        %v1 = queue.read(%k1)
        %v2 = addf(%v1, %v2)
        queue.write(%v2, %m2)
        queue.return
    }
    %s4 = queue.memcpy(%s3, %arg2, %2, %dma)
    %s5 = queue.await(%s4)
    queue.return}
```

Rewriting Aten Dialect

```
func @graph( %arg0: memref<16xf32>, %arg1:
    memref<16xf32>) -> memref<16xf32> {

    %m0 = aten.acap_alloc()
    %s0 = aten.acap_copy(%m0, %arg0)
    %v0 = aten.acap_tensor_load(%m0, %s0)

    %m1 = aten.acap_alloc()
    %s1 = aten.acap_copy(%m1, %arg0)
    %v1 = aten.acap_tensor_load(%m1, %s1)


    %s2 = aten.constant()
    %m2 = aten.acap_alloc()
    %v2 = aten.add(%v0, %v1, %s2)
    %s3 = aten.acap_tensor_store(%v2, %m2)
    %m3 = aten.acap_alloc()
    %s4 = aten.acap_copy(%m2, %m3, %s3)
    %s5 = aten.acap_wait_all(%s4)
    return %m3
}
```

```
queue.launch(%kernel, %mem, %dma, %arg0, %arg1, %arg2 =
    %aie, %SRAM, %DMA, %0, %1, %2) in (%s2, %microblaze){
    %start = queue.control_start ()
    %m0 = queue.alloc %mem, [16], f32
    %m1 = queue.alloc %mem, [16], f32
    %m2 = queue.alloc %mem, [16], f32
    %s0 = queue.memcpy(%start, %arg0, %m0, %dma)
    %s1 = queue.memcpy(%start, %arg1, %m1, %dma)
    %s2 = queue.control_and(%s0, %s1)
    %s3 = queue.launch(%k0, %k1, %k2 = %m0, %m1, %m2) in
    (%s2, %kernel) {
        %v0 = queue.read(%k0)
        %v1 = queue.read(%k1)
        %v2 = addf(%v1, %v2)
        queue.write(%v2, %m2)
        queue.return
    }
    %s4 = queue.memcpy(%s3, %arg2, %2, %dma)
    %s5 = queue.await(%s4)
    queue.return}
```

%SRAM = create_mem [1024], f32, SRAM

launch(%mem = %SRAM)

Rewriting Aten Dialect



```
queue.launch(%kernel, %mem, %dma, %arg0, %arg1, %arg2 =
%aie, %SRAM, %DMA, %0, %1, %2) in (%s2, %microblaze){
    %start = queue.control_start ()
    %m0 = queue.alloc %mem, [16], f32
    %m1 = queue.alloc %mem, [16], f32
    %m2 = queue.alloc %mem, [16], f32
    %s0 = queue.memcpy(%start, %arg0, %1, %dma)
    %s1 = queue.memcpy(%start, %arg1, %4, %dma)
    %s2 = queue.control_and(%s0, %s1)
    %s3= queue.launch(%k0, %k1, %k2 = %m0, %m1, %m2) in
    (%s2, %kernel) {
        %v0 = queue.read(%k0)
        %v1 = queue.read(%k1)
        %v2 = addf(%v1, %v2)
        queue.write(%v2, %m2)
        queue.return
    }
    %s4 = queue.memcpy(%s3, %arg2, %2, %dma)
    %s5 = queue.await(%s4)
    queue.return}
```

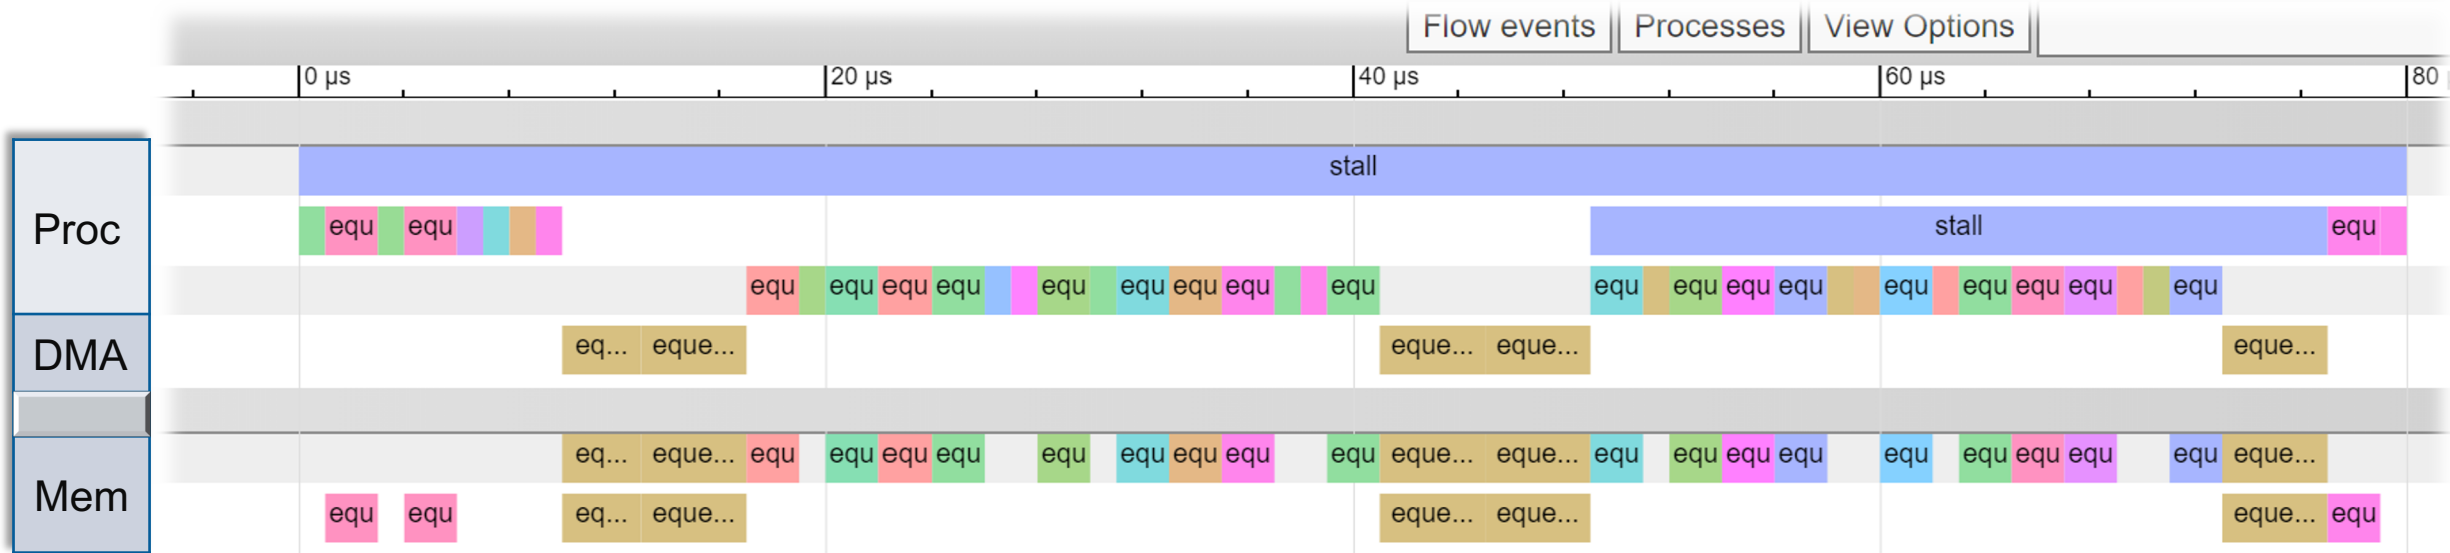


```

%start_signal = "equeue.control_start"():()->!equeue.signal
%done_signal, %result = equeue.launch (%act, %weight, %aie_core_local, %dma, %m0, %m1 = %act_in, %weight_in,
%aie_core, %accel_dma, %sram, %aie_mem : tensor<16xf32>, tensor<5xf32>, i32, i32, i32, i32)
in (%start_signal, %accel_core) : tensor<12xf32> {
    %weight_sram = equeue.alloc %m0, [5], f32 : !equeue.container<tensor<5xf32>, i32>
    "equeue.write"(%weight, %weight_sram): (tensor<5xf32>, !equeue.container<tensor<5xf32>, i32>)->()
    %act_sram = equeue.alloc %m0, [16], f32: !equeue.container<tensor<16xf32>, i32>
    "equeue.write"(%act, %act_sram): (tensor<16xf32>, !equeue.container<tensor<16xf32>, i32>)->()
    %output_sram = equeue.alloc %m0, [16], f32 : !equeue.container<tensor<12xf32>, i32>
    %weight_mem = equeue.alloc %m1, [5], f32:!equeue.container<tensor<5xf32>, i32>
    %act_mem = equeue.alloc %m1, [5], f32 :!equeue.container<tensor<5xf32>, i32>
    %ofmap_mem = equeue.alloc %m1, [1], f32 :!equeue.container<f32, i32>
    %start_memcpy = "equeue.control_start"():()->!equeue.signal
    %dma_done0 = "equeue.memcpy"(%start_memcpy, %weight_sram, %weight_mem, %dma): (!equeue.signal, !equeue.container<tensor<5xf32>,i32>, !equeue.container<tensor<5xf32>,i32>, i32) -> !equeue.signal
    %c0 = constant 0:index
    %c12 = constant 2:index
    %c1 = constant 1:index
    %compute_done = scf.for %k = %c0 to %c12 step %c1 iter_args(%dma_start = %start_memcpy) -> (!equeue.signal) {
        %dma_done1 = "equeue.memcpy"(%dma_start, %act_sram, %act_mem, %dma, %k): (!equeue.signal, !equeue.container<tensor<16xf32>,i32>, !equeue.container<tensor<5xf32>, i32>, i32, index)
        %start_compute = "equeue.control_and"(%dma_done0, %dma_done1):(!equeue.signal, !equeue.signal)->!equeue.signal
        %compute_once = equeue.launch (%act_mem_local, %weight_mem_local, %ofmap_mem_local, %offset =
        %act_mem, %weight_mem, %ofmap_mem, %k: !equeue.container<tensor<5xf32>, i32>, !equeue.container<tensor<5xf32>, i32>,
        !equeue.container<f32, i32>, index)
        in (%start_compute, %aie_core_local) {
            %const0 = constant 0.0:f32
            "equeue.write"(%const0, %ofmap_mem_local): (f32, !equeue.container<f32, i32>)->()
            %cst0 = constant 0:index
            %cst1 = constant 1:index
            %cst5 = constant 2:index
            scf.for %i = %cst0 to %cst5 step %cst1 {
                %j = affine.apply affine_map<(d0,d1)->(d0+d1)>(%offset,%i)
                %ifmap = "equeue.read" (%act_mem_local,%j):(!equeue.container<tensor<5xf32>, i32>, index)->f32
                %filter = "equeue.read" (%weight_mem_local,%i):(!equeue.container<tensor<5xf32>, i32>, index)->f32
                %ofmap = "equeue.read" (%ofmap_mem_local):(!equeue.container<f32, i32>) -> f32
                %psum = mulf %filter, %ifmap: f32
                %ofmap_flight = addf %ofmap, %psum: f32
                "equeue.write"( %ofmap_flight, %ofmap_mem_local):(f32, !equeue.container<f32, i32>)->()
                "scf.yield"():()->() }
            "equeue.return"():()->() }
        %ofmap_done = "equeue.memcpy"(%compute_once, %ofmap_mem, %output_sram, %dma, %k):(!equeue.signal, !equeue.container<f32, i32>, !equeue.container<tensor<12xf32>, i32>, i32, index)
        "scf.yield"(%ofmap_done):(!equeue.signal)->() }
    "equeue.await"(%compute_done):(!equeue.signal)->()
    %act_out = "equeue.read" (%output_sram):(!equeue.container<tensor<12xf32>, i32> ) -> tensor<5xf32>
    equeue.dealloc %act_sram, %weight_sram, %output_sram, %act_mem, %weight_mem, %ofmap_mem: !equeue.container<tensor<16xf32>, i32>, !equeue.container<tensor<5xf32>, i32>, !equeue.container<tensor<12xf32>
    "equeue.await"(%done_signal):(!equeue.signal)->()
    return %result: tensor<12xf32>

```






Visualized Simulation



- ▶ Concurrency
- ▶ Synchronized scheduling
- ▶ Model execution time
- ▶ Await


Case Study – Input Stationary on Systolic Array


Input stationary



a2	b2	c2	d2		
e2	f2	a1	b1	c1	d1
i2	j2	e1	f1	g1	h1
m2	n2	i1	j1	k1	l1
		m1	n1	o1	p1

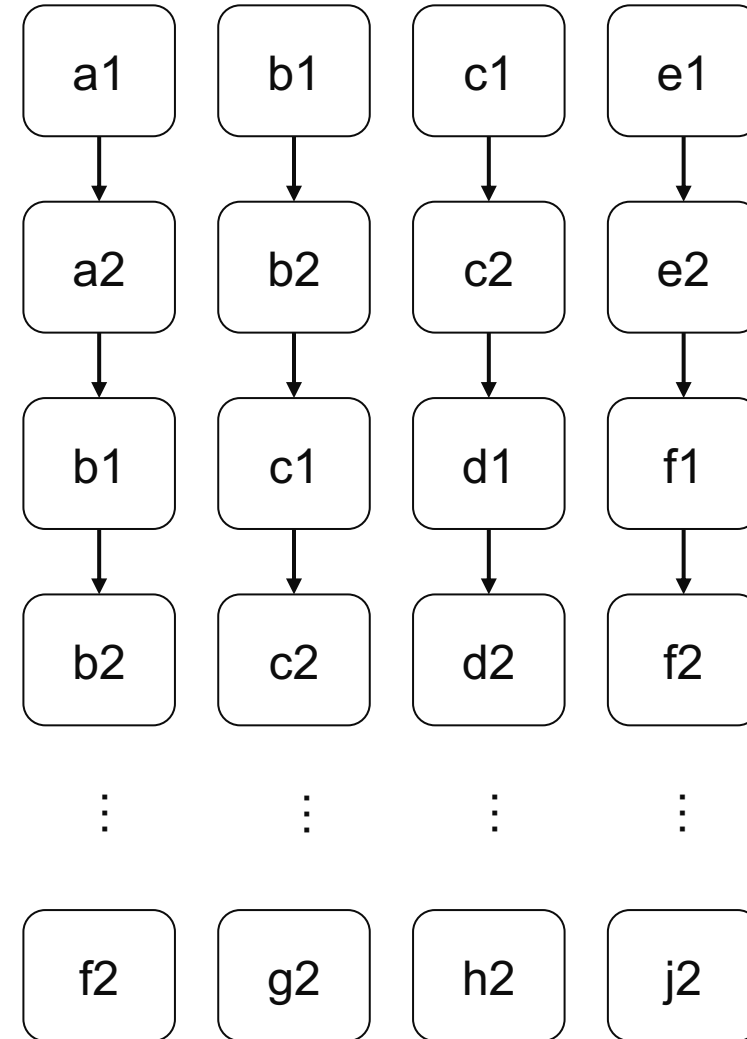
a1	b1	c1	e1
a2	b2	c2	e2
b1	c1	d1	f1
b2	c2	d2	f2
e1	f1	g1	i1
e2	f2	g2	i2
f1	g1	h1	j1
f2	g2	h2	j2














Read inputs





a2	b2	c2	d2		
e2	f2	a1	b1	c1	d1
i2	j2	e1	f1	g1	h1
m2	n2	i1	j1	k1	l1
		m1	n1	o1	p1

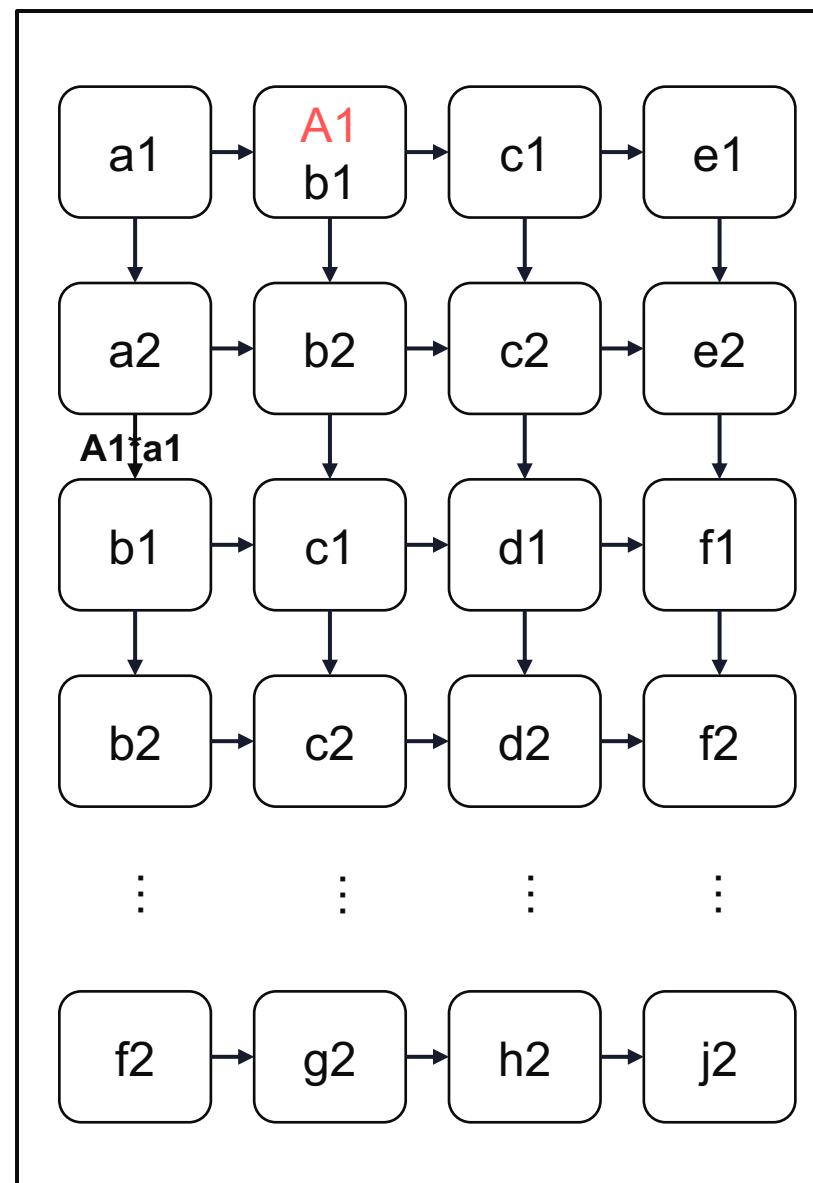





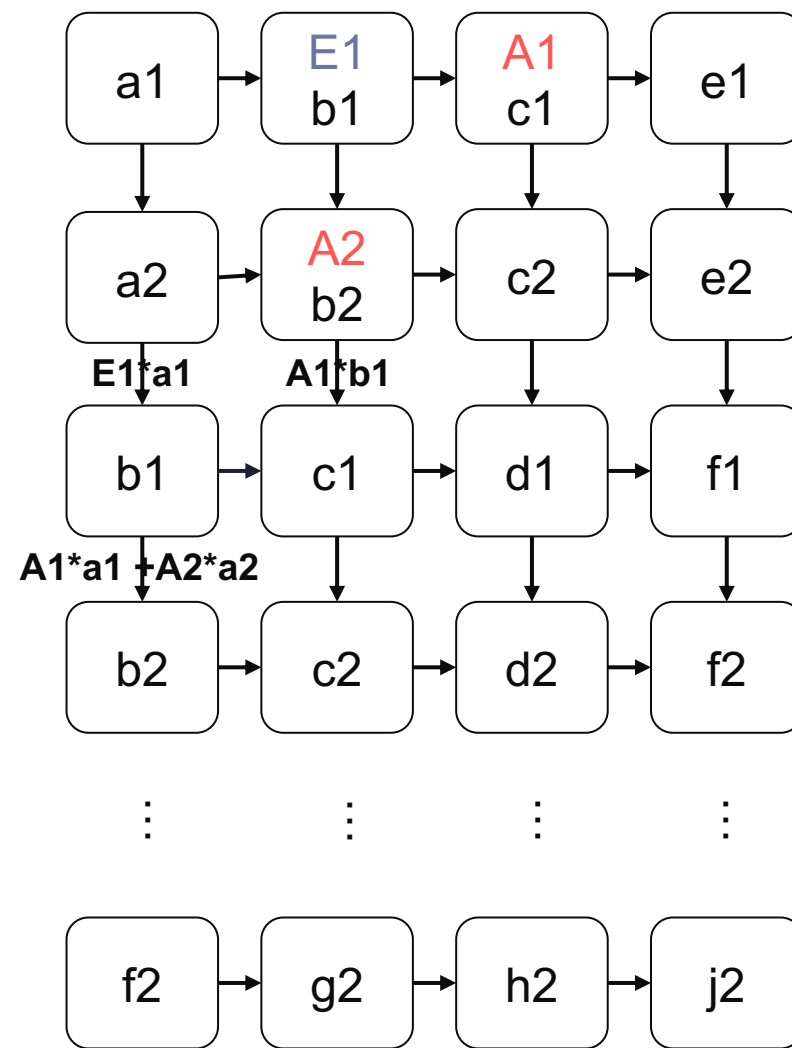
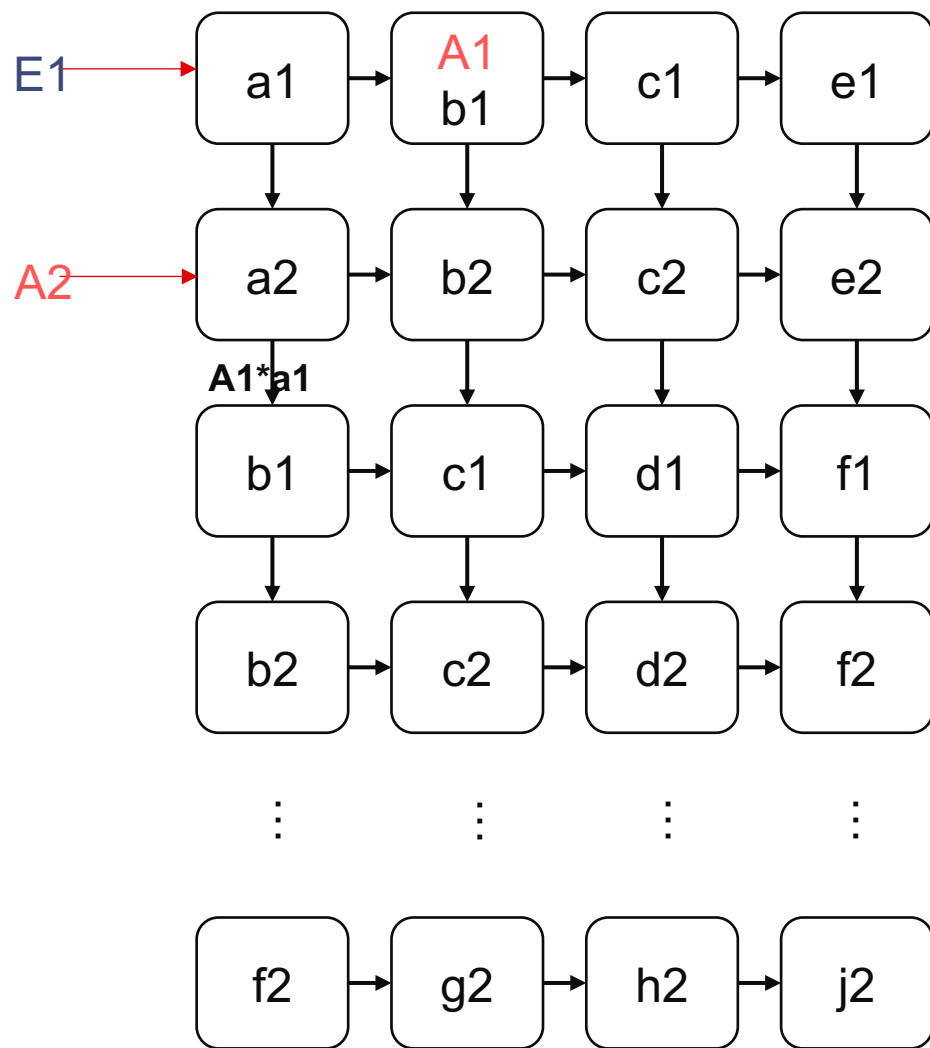


Compute




compute









Example Generator

- ▶ Generator create large example (mlir) with limited lines of codes (c++)
 - For loop and events
 - **declarative builders API** (<https://mlir.llvm.org/docs/EDSC/>)
- ▶ Generator
 - https://github.com/cucapra/event_queue/blob/scalable/lib/Generator/EQueueScaleSim.cpp
- ▶ MLIR Example
 - https://github.com/cucapra/event_queue/blob/scalable/test/EQueue/scale_sim_os.mlir
- ▶ Visualizer
 - <chrome://tracing/>

Representing Parallel Flow in Equeue Dialect

```
// start signal
```

```
start = equeue.start()
```

```
// parallel for
```

```
for h in arr_height:
```

```
    done, ... = equeue.launch(...) in (start, core[h]) {...}
```

```
    if h = 0:
```

```
        prev_done = done
```

```
    else:
```

```
        prev_done = equeue.control_and(done, prev_done)
```

```
// await for done signal to proceed
```

```
equeue.await(prev_done)
```

Input Stationary Workflow

// Moving in inputs

for t in 0...filter_w * filter_h* channel:

 //inside each pe, read out current map

 //shift current map

 //memcpy from sram to pe registers

// compute and output results

for t in 0...output_size + arr_h + arr_w:

 // memcpy weights from sram to pe registers

 // read the current values in pe registers and calculate ofmap

 // update pe registers and potentially write output to sram

Moving in inputs

// Moving in inputs

for t in 0...filter_w * filter_h * channel:

// inside each pe, read out current map

start = equeue.start() //parallel

for h in arr_height:

for w in arr_width:

done, ifmap_flight[h][w] = equeue.launch(ifmap_buffer=pe[h][w].ifmap_buffer) in (start, pe[h][w].core) {

ifmap = equeue.read(ifmap_buffer)

equeue.yield ifmap

}

if h= 0 and w = 0:

prev_done = done

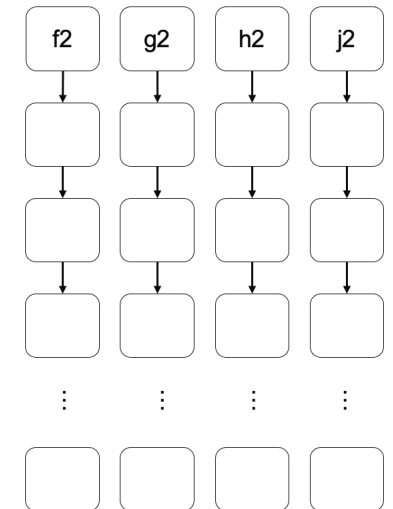


else:

prev_done = equeue.control_and(done, prev_done)

equeue.await(prev_done)

// shift current map

// memcpy from sram to pe registers



Moving in inputs

// Moving in inputs

```
for t in 0...filter_w * filter_h* channel:
```

```
    // inside each pe, read out current map
```

// shift current map

```
start = equeue.start() //parallel
```

```
for h in 1...arr_height:
```

```
    for w in arr_width:
```

```
        done = equeue.launch(ifmap_buffer=pe[h][w].ifmap_buffer, ifmap_flight=ifmap_flight[h-1][w]) in (start, pe[h][w].core) {  
            equeue.write(ifmap_flight, ifmap_buffer)
```

```
        }
```

```
        if h= 1 and w = 0:
```


```
            prev_done = done
```

```
        else:
```

```
            prev_done = equeue.control_and(done, prev_done)
```

```
equeue.await(prev_done)
```

```
// memcpy from sram to pe registers
```



Moving in inputs

// Moving in inputs

for t in 0...filter_w * filter_h* channel:

// inside each pe, read out current map

// shift current map

// memcpy from sram to pe registers

start = equeue.start()

for w in arr_width:

//different bank to ensure parallel execution

done = equeue.memcopy(start, ifmap_sram, ifmap_buffer=pe[0][w].ifmap_buffer, offset, bank)




if w = 0:

prev_done = done

else:

prev_done = equeue.control_and(done, prev_done)

equeue.await(prev_done)



Compute and Output Results

```
// compute and output results
```

```
for t in 0...output_size + arr_h + arr_w:
```

```
    // memcpy weights from sram to pe registers
```

```
    start = equeue.start()
```

```
    for n in 0...weights_num (step=write_direction):
```

```
        done = equeue.memcopy(start, weights_sram, weight_buffer=pe[n][0].weight_buffer, offset, bank)
```

```
        if w = 0:
```

```
            prev_done = done
```


```
        else:
```

```
            prev_done = equeue.control_and(done, prev_done)
```

```
    equeue.await(prev_done)
```

```
// read the current values in pe registers and calculate ofmap
```

```
// update pe registers and potentially write output to sram
```



Compute and Output Results

```
// compute and output results
```

```
for t in 0...output_size + arr_h + arr_w:
```

```
// memcopy weights from sram to pe registers
```

```
// read the current values in pe registers and calculate ofmap
```

```
start = equeue.start()
```

```
for h in arr_height:
```

```
for w in arr_width:
```

```
done, weight_flight[h][w], ofmap_flight[h][w]= equeue.launch( start, ifmap_buffer=pe[h][w].ifmap_buffer,
```

```
ofmap_buffer=pe[h][w].ofmap_buffer,
```

```
weight_buffer=pe[h][w].weight_buffer
```

```
) in (start, pe[h][w].core) {
```

```
ifmap = ifmap_buffer
```

```
weight = weight_buffer
```

```
ofmap_old = equeue.read(ofmap_buffer)
```

$$\text{ofmap} = \text{ifmap} * \text{weight} + \text{ofmap_old}$$

enqueue.yield weights, ofmap


}

```
if h= 0 and w = 0: prev_done = done
```

```
else: prev_done = equeue.control_and(done, prev_done)
```

```
enqueue.await(prev_done)
```

```
// update pe registers and potentially write output to sram
```



Compute and Output Results

```
// compute and output results
```

```
for t in 0...output_size + arr_h + arr_w:
```

```
    // memcpy weights from sram to pe registers
```

```
    // read the current values in pe registers and calculate ofmap
```

```
    // update pe registers and potentially write output to sram
```

```
    start = equeue.start()
```

```
    for r in arr_height:
```

```
        for c in arr_width:
```

```
            done = equeue.launch(
```

```
                start,
```

```
                wflight=wflight[r][c],
```

```
                oflight=oflight[r][c],
```

```
                wbuffer=wbuffer2s[r][c+1],
```

```
                obuffer=obuffer2s[r+1][c],
```

```
            ){
```

```
                //no outputs
```

```
                equeue.write(oflight, obuffer)
```

```
                equeue.write(wflight, wbuffer)
```


```
                if r+c <= t and t <= ofmap_size + arr_h + c:
```

```
                    equeue.writeop(oflight, out_sram)
```

```
            }
```

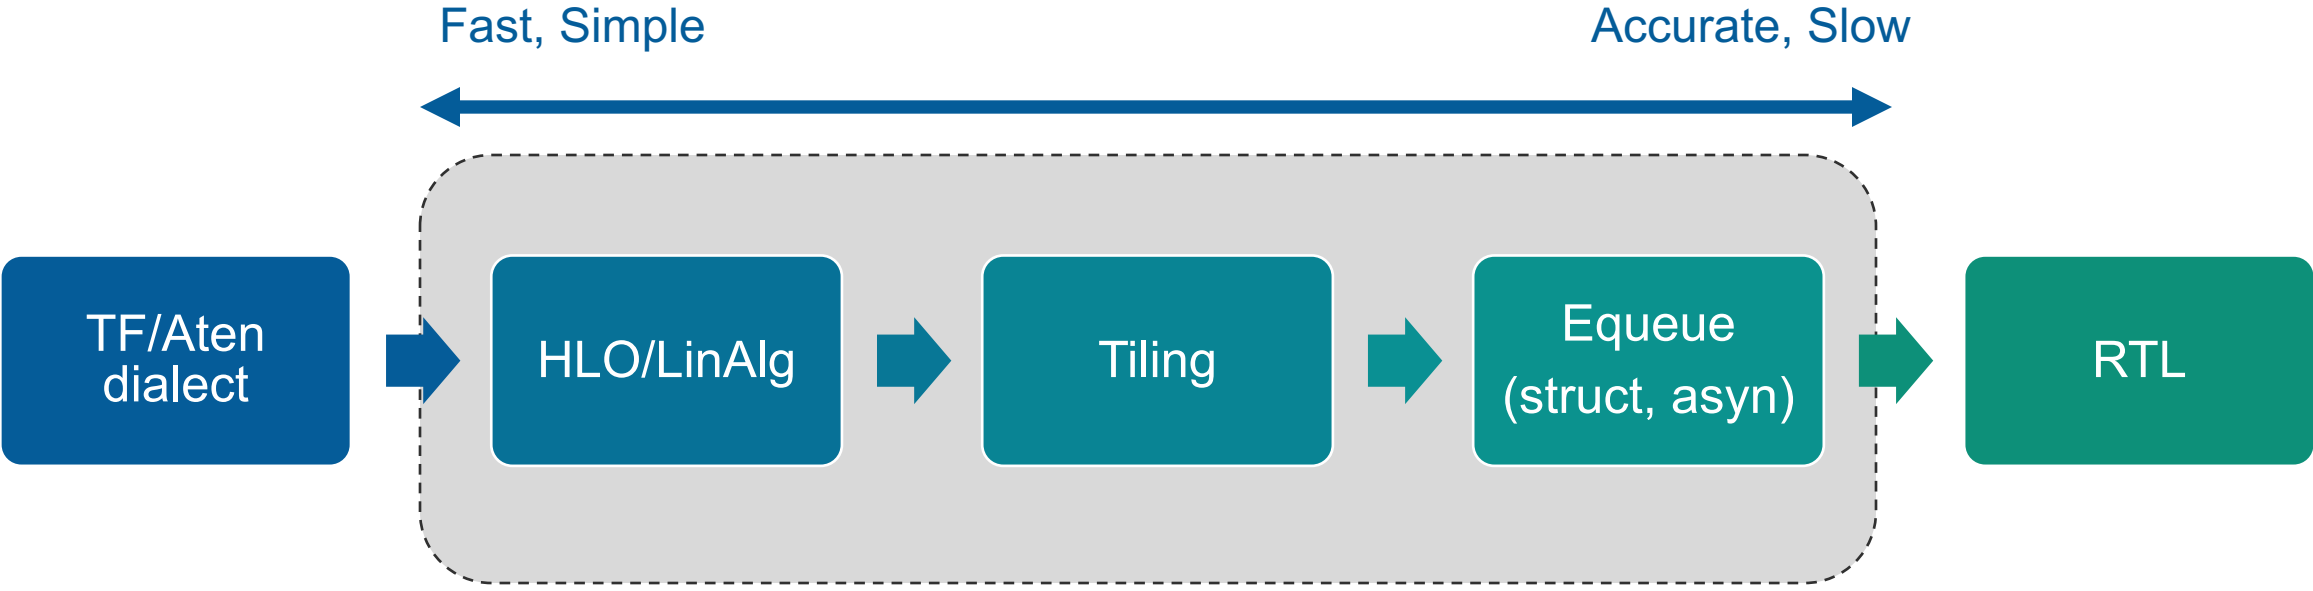
```
        }
```

```
    equeue.await(prev_done)
```



High Level Abstractions

Lowering Pipeline

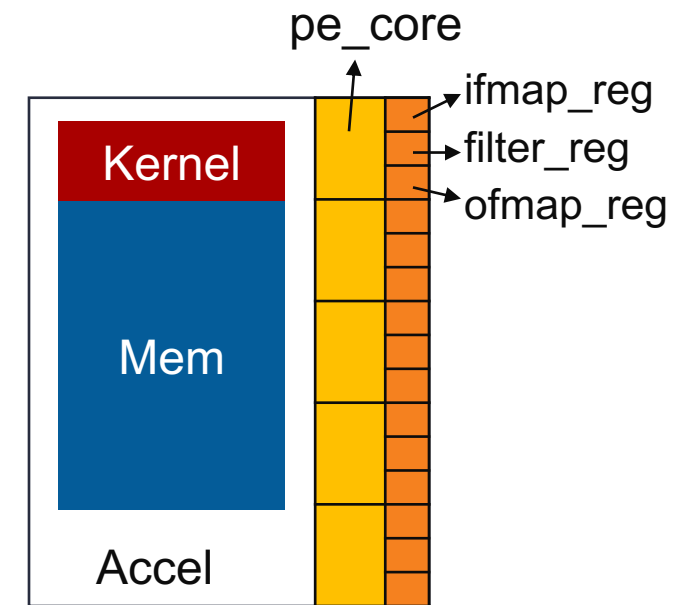


Different Level of Abstractions in EQueue Dialect

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W)  
}
```


Different Level of Abstractions in EQueue Dialect

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W)  
}
```



Different Level of Abstractions in EQueue Dialect

Tiling Strategy

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W)  
}
```

reg_size, etc

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W/(Ah*Aw), Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/(Ah*Aw, W, Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/Ah, W/*Aw, Ah, Aw)  
}
```

Different Level of Abstractions in EQueue Dialect

Tiling Strategy

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W/(Ah*Aw), Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/(Ah*Aw, W, Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/Ah, W/*Aw, Ah, Aw)  
}
```

Data Movement

Weight Stationary

Row Stationary

Output Stationary

Different Level of Abstractions in EQueue Dialect

Tiling Strategy

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H, W/(Ah*Aw), Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/(Ah*Aw, W, Ah, Aw)  
}
```

```
launch (...) in (%start, %accel) {  
  generic (E, F, H/Ah, W/*Aw, Ah, Aw)  
}
```

Data Movement

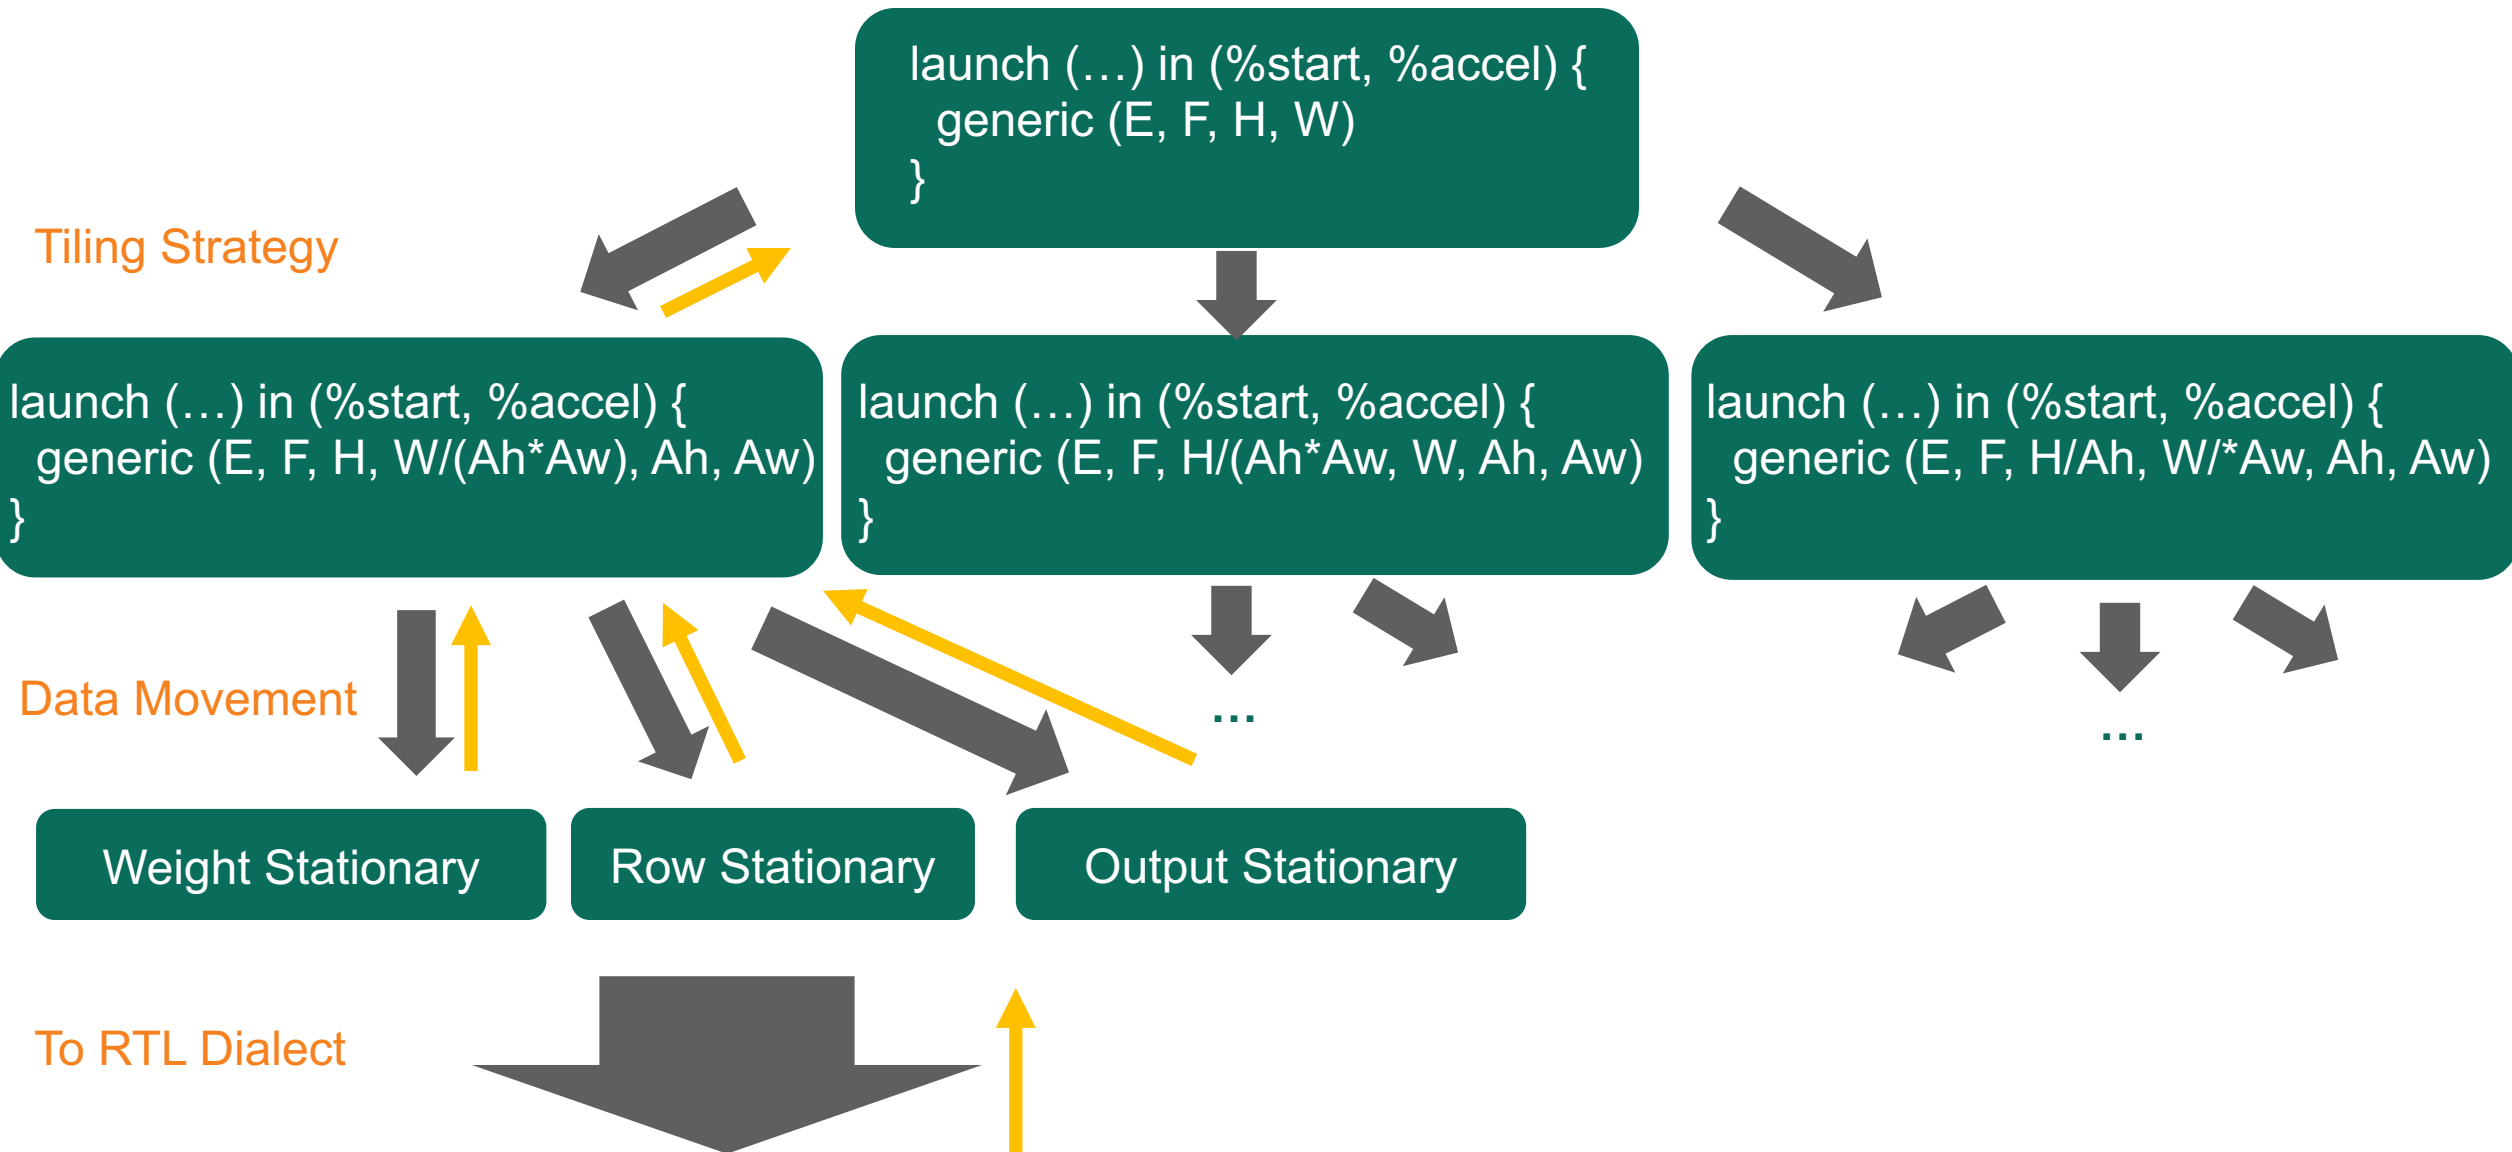
Weight Stationary

Row Stationary

Output Stationary

To RTL Dialect

Different Levels of Abstractions in EQueue Dialect



Background: Linalg Dialect

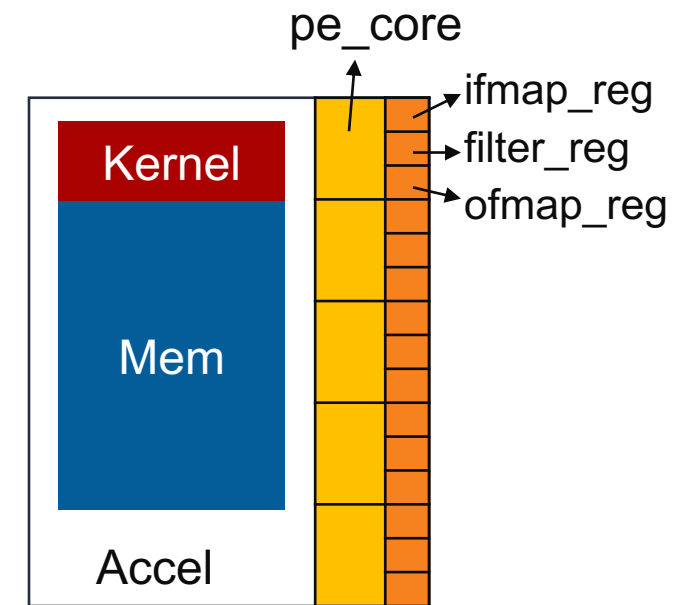
```
#accesses = [  
    affine_map<(E, F, H, W) -> (E, F)>,  
    affine_map<(E, F, H, W) -> (H, W)>,  
    affine_map<(E, F, H, W) -> (E, F)>  
]  
#trait = {  
    args_in = 2,  
    args_out = 1,  
    iterator_types = ["reduction", "reduction", "parallel", "parallel"],  
    indexing_maps = #accesses  
}  
func @compute(%A: memref<3x3xf32>, %B: memref<5x5xf32>, %C:  
memref<3x3xf32>) {  
    linalg.generic #trait %A, %B, %C {  
        ^bb0(%a: f32, %b: f32, %c: f32):  
            %d = "some_compute"(%a, %b): (f32, f32)->f32  
            linalg.yield %d: f32  
    } : memref<3x3xf32>, memref<5x5xf32>, memref<3x3xf32>  
    return  
}
```

Background: Representation of Linalg in Loops

```
func @compute(%arg0: memref<3x3xf32>, %arg1: memref<5x5xf32>, %arg2: memref<3x3xf32>) {  
  %c3 = constant 3 : index  
  %c5 = constant 5 : index  
  %c0 = constant 0 : index  
  %c1 = constant 1 : index  
  scf.for %arg3 = %c0 to %c3 step %c1 {  
    scf.for %arg4 = %c0 to %c3 step %c1 {  
      scf.parallel (%arg5, %arg6) = (%c0, %c0) to (%c5, %c5) step %c1 {  
        %0 = load %arg0[%arg3, %arg4] : memref<3x3xf32>  
        %1 = load %arg1[%arg5, %arg6] : memref<5x5xf32>  
        %2 = "some_compute"(%0, %1) : (f32, f32) -> f32  
        store %2, %arg2[%arg3, %arg4] : memref<3x3xf32>  
        scf.yield  
      }  
    }  
  }  
  return  
}
```

Equeue - Highest Level

```
#accesses = [  
  affine_map<(E, F, H, W) -> (E, F)>,  
  affine_map<(E, F, H, W) -> (H, W)>,  
  affine_map<(E, F, H, W) -> (E, F)>  
]  
#trait = {  
  args_in = 2,  
  args_out = 1,  
  iterator_types = ["reduction", "reduction", "parallel", "parallel"],  
  indexing_maps = #accesses  
}  
equeue.launch(%A=%ifmap, %B = %filter, %C = %ofmap) in (%start_signal, %accel_kernel){  
  linalg.generic #trait %A, %B, %C {  
    ^bb0(%a: f32, %b: f32, %c: f32):  
      %d = "some_compute"(%a, %b): (f32, f32)->f32  
      linalg.yield %d: f32  
  } : memref<3x3xf32>, memref<5x5xf32>, memref<3x3xf32>  
  equeue.return  
}
```



Tiling Linalg

```
equeue.launch(%arg0=%ifmap, %arg1 = %filter,  
%arg2 = %ofmap) in (%start_signal, %accel_kernel){  
  ... // constants  
  scf.for %arg3 = %c0 to %c3 step %c1 {  
    scf.for %arg4 = %c0 to %c3 step %c1 {  
      scf.parallel (%arg5) = (%c0) to (%c5) step (%c1) {  
        %0 = affine.min #map0(%arg3)  
        %1 = affine.min #map0(%arg4)  
        %2 = subview %arg0[%arg3, %arg4] [%0, %1] [1, 1] : memref<3x3xf32> to memref<?x?xf32, #map1>  
        %3 = affine.min #map2(%arg5)  
        %4 = subview %arg1[0, %arg5] [5, %3] [1, 1] : memref<5x5xf32> to memref<5x?xf32, #map3>  
        %5 = affine.min #map0(%arg3)  
        %6 = affine.min #map0(%arg4)  
        %7 = subview %arg2[%arg3, %arg4] [%5, %6] [1, 1] : memref<3x3xf32> to memref<?x?xf32, #map1>  
        linalg.generic {args_in = 2 : i64, args_out = 1 : i64, indexing_maps = [#map4, #map5, #map4], iterator_types =  
["reduction", "reduction", "parallel", "parallel"]} %2, %4, %7 {  
          ^bb0(%arg6: f32, %arg7: f32, %arg8: f32): // no predecessors  
            %8 = "some_compute"(%arg6, %arg7) : (f32, f32) -> f32  
            linalg.yield %8 : f32  
          }: memref<?x?xf32, #map1>, memref<5x?xf32, #map3>, memref<?x?xf32, #map1> }}}  
    equeue.return }
```

```
#map0 = affine_map<(d0) -> (1, -d0 + 3)>  
#map1 = affine_map<(d0, d1)[s0] -> (d0 * 3 + s0 + d1)>  
#map2 = affine_map<(d0) -> (1, -d0 + 5)>  
#map3 = affine_map<(d0, d1)[s0] -> (d0 * 5 + s0 + d1)>  
#map4 = affine_map<(d0, d1, d2, d3) -> (d0, d1)>  
#map5 = affine_map<(d0, d1, d2, d3) -> (d2, d3)>
```

Tiling Linalg

```
equeue.launch(%arg0=%ifmap, %arg1 = %filter,  
%arg2 = %ofmap) in (%start_signal, %accel_kernel){  
  ... // constants  
  scf.for %arg3 = %c0 to %c3 step %c1 {  
    scf.for %arg4 = %c0 to %c3 step %c1 {  
      scf.parallel (%arg5) = (%c0) to (%c5) step (%c1) {  
  
        ...  
  
      }: memref<?x?xf32, #map1>, memref<5x?xf32, #map3>, memref<?x?xf32, #map1> }}}  
equeue.return }
```

```
#map0 = affine_map<(d0) -> (1, -d0 + 3)>  
#map1 = affine_map<(d0, d1)[s0] -> (d0 * 3 + s0 + d1)>  
#map2 = affine_map<(d0) -> (1, -d0 + 5)>  
#map3 = affine_map<(d0, d1)[s0] -> (d0 * 5 + s0 + d1)>  
#map4 = affine_map<(d0, d1, d2, d3) -> (d0, d1)>  
#map5 = affine_map<(d0, d1, d2, d3) -> (d2, d3)>
```

Parallel Loop Requires Control-Structure Mapping!

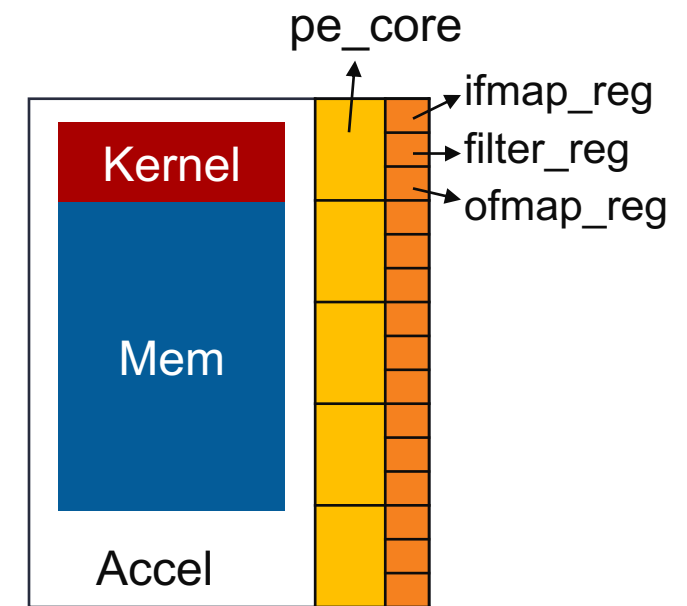
```

equeue.launch(%arg0=%ifmap, %arg1 = %filter, %arg2 = %ofmap) in (%start_signal,
%accel_kernel){
  ... // constants
  scf.for %arg3 = %c0 to %c3 step %c1 {
    scf.for %arg4 = %c0 to %c3 step %c1 {
      %start_pe = equeue.control_start()
      equeue.launch(%local_ifmap=%pe0_ifmap, %local_filter=%pe0_filter,
%local_ofmap=%pe0_ofmap) in (%start_pe, %pe0_core){
        linalg.generic {args_in = 2 : i64, args_out = 1 : i64, indexing_maps = [#map4,
#map5, #map4], iterator_types = ["reduction", "reduction", "parallel", "parallel"]}
        %local_ifmap, %local_filter, %local_ofmap {
          ^bb0(%arg6: f32, %arg7: f32, %arg8: f32): // no predecessors
            %8 = "some_compute"(%arg6, %arg7) : (f32, f32) -> f32
            linalg.yield %8 : f32
          }: memref<?x?xf32, #map1>, memref<5x?xf32, #map3>, memref<?x?xf32,
#map1> }

        equeue.launch(%local_ifmap=%pe1_ifmap, %local_filter=%pe1_filter,
%local_ofmap=%pe1_ofmap) in (%start_pe, %pe1_core){...}

        ... //pe2, pe3, pe4
      }}
    equeue.return }
  }
}

```



Parallel Loops to Equeue Structure – Data Movements

```
equeue.launch(%arg0=%ifmap, %arg1 = %filter, %arg2 = %ofmap) in (%start_signal, %accel_kernel){  
    ... // constants  
    scf.for %arg3 = %c0 to %c3 step %c1 {  
        scf.for %arg4 = %c0 to %c3 step %c1 {  
            %start_pe = enqueue.control_start()  
            %done_copy0 = enqueue.memcpy(%start_pe, %2, %local_ifmap,  
                                         %offset0, %bank0)  
            %done_copy1 = enqueue.memcpy(%start_pe, %2, %local_ifmap,  
                                         %offset0, %bank0)  
            ...  
  
            enqueue.launch(%local_ifmap=%pe0_ifmap, %local_filter=%pe0_filter,  
                           %local_ofmap=%pe0_ofmap) in (%done_copy0, %pe0_core){...}  
            enqueue.launch(%local_ifmap=%pe1_ifmap, %local_filter=%pe1_filter,  
                           %local_ofmap=%pe1_ofmap) in (%done_copy1, %pe1_core){...}  
            ... //pe1, .. pe2, pe3, pe4  
        }  
    }  
    enqueue.return }
```

Parallel Loops to Equeue Structure – Weight Stationary

```
equeue.launch(%arg0=%ifmap, %arg1 = %filter, %arg2 = %ofmap) in (%start_signal, %accel_kernel){
  ... // constants
  scf.for %arg3 = %c0 to %c3 step %c1 {
    scf.for %arg4 = %c0 to %c3 step %c1 {
      %start_pe = equeue.control_start()
      %done_copy0 = equeue.memcpy(%start_pe, %2, %local_ifmap,
                                %offset0, %bank0)
      %done_copy1 = equeue.memcpy(%start_pe, %2, %local_ifmap,
                                %offset0, %bank0)

      ...
      %done_copy = equeue.await(%done_copy0, %done_copy1, ...)
      scf.for %arg4 = %c0 to %c5 step %c1 {
        equeue.launch(%local_ifmap=%pe0_ifmap, %local_filter=%pe0_filter,
                      %local_ofmap=%pe0_ofmap) in (%done_copy, %pe0_core){...}
        equeue.launch(%local_ifmap=%pe1_ifmap, %local_filter=%pe1_filter,
                      %local_ofmap=%pe1_ofmap) in (%done_copy, %pe1_core){...}
        ... //pe1, .. pe2, pe3, pe4
      }
    }
  }
  equeue.return }
```

Ongoing Work & Future Directions

- ▶ Different levels of abstractions
- ▶ Generate fast simulation
- ▶ Functional correctness verification
- ▶ Representing dynamic execution
 - Sparsity!

Thank you!